

Distributed Coordination

Luís Osório

■ The Failure

- The computer system failure (CPU, memory, power supply, etc)
 - Commonly total-failure
- Software running error, e.g., memory violation, communication error, etc. leading to process/thread failure
 - Commonly partial-failure
- In ([link](#), or the [pdf](#)) Schneider summarizes definitions, from both Lamport ([link](#)), the Byzantine failure, and, from himself ([link](#)) Fail-stop failure:
 - “**Byzantine Failures**. The component can exhibit arbitrary and malicious behavior, perhaps involving collusion with other faulty components [Lamport et al 82].”
 - “**Fail-stop Failures**. In response to a failure, the component changes to a state that permits other components to detect that a failure has occurred and then stops [Schneider 84].”
- Byzantine failures as discussed by Lamport ([link](#))
 - “A **reliable computer system** must be able to cope with the failure of one or more of its components. A failed component may exhibit a type of behavior that is often overlooked--namely, sending conflicting information to different parts of the system. The problem of coping with this type of failure is expressed abstractly as the **Byzantine Generals Problem**.”


The Problem Motivating Distributed Coordination

■ Consider a Service element responsible for executing a Finite State Machine

- The next state depends on inputs and current state
- State transition occurs when a special input, the event clock is triggered (synchronous)

■ As discussed by Schneider in ([link](#))

- “... every **protocol** we know of that **employs replication** - be it for masking failures or simply to facilitate cooperation without centralized control - can be **derived** using the **state machine** approach.”
- The question is:
 - How to guarantee an FSM evolves consistently independent of a failure that may occur?
- Consider, a reliable ServiceFSM.





From: Zookeeper, Distributed Process Coordination;
ISBN: 9781449361303, O'Reilly, 2013

Fault Tolerance with the Zookeeper System

Distributed Coordination

■ Apache Curator

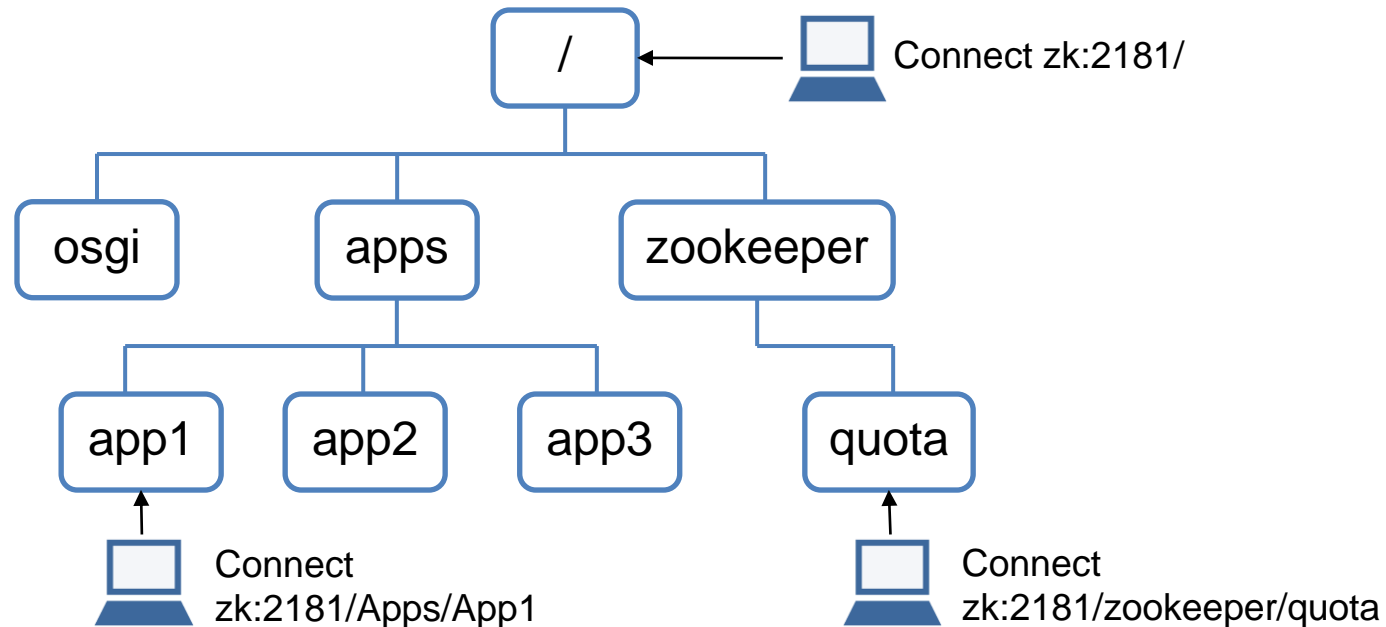
- Enhanced Client library for Zookeeper
 - <https://curator.apache.org/index.html>
- References
 - <https://www.baeldung.com/apache-curator>
- Source and Examples
 - <https://github.com/apache/curator>

■ Spring Cloud Zookeeper

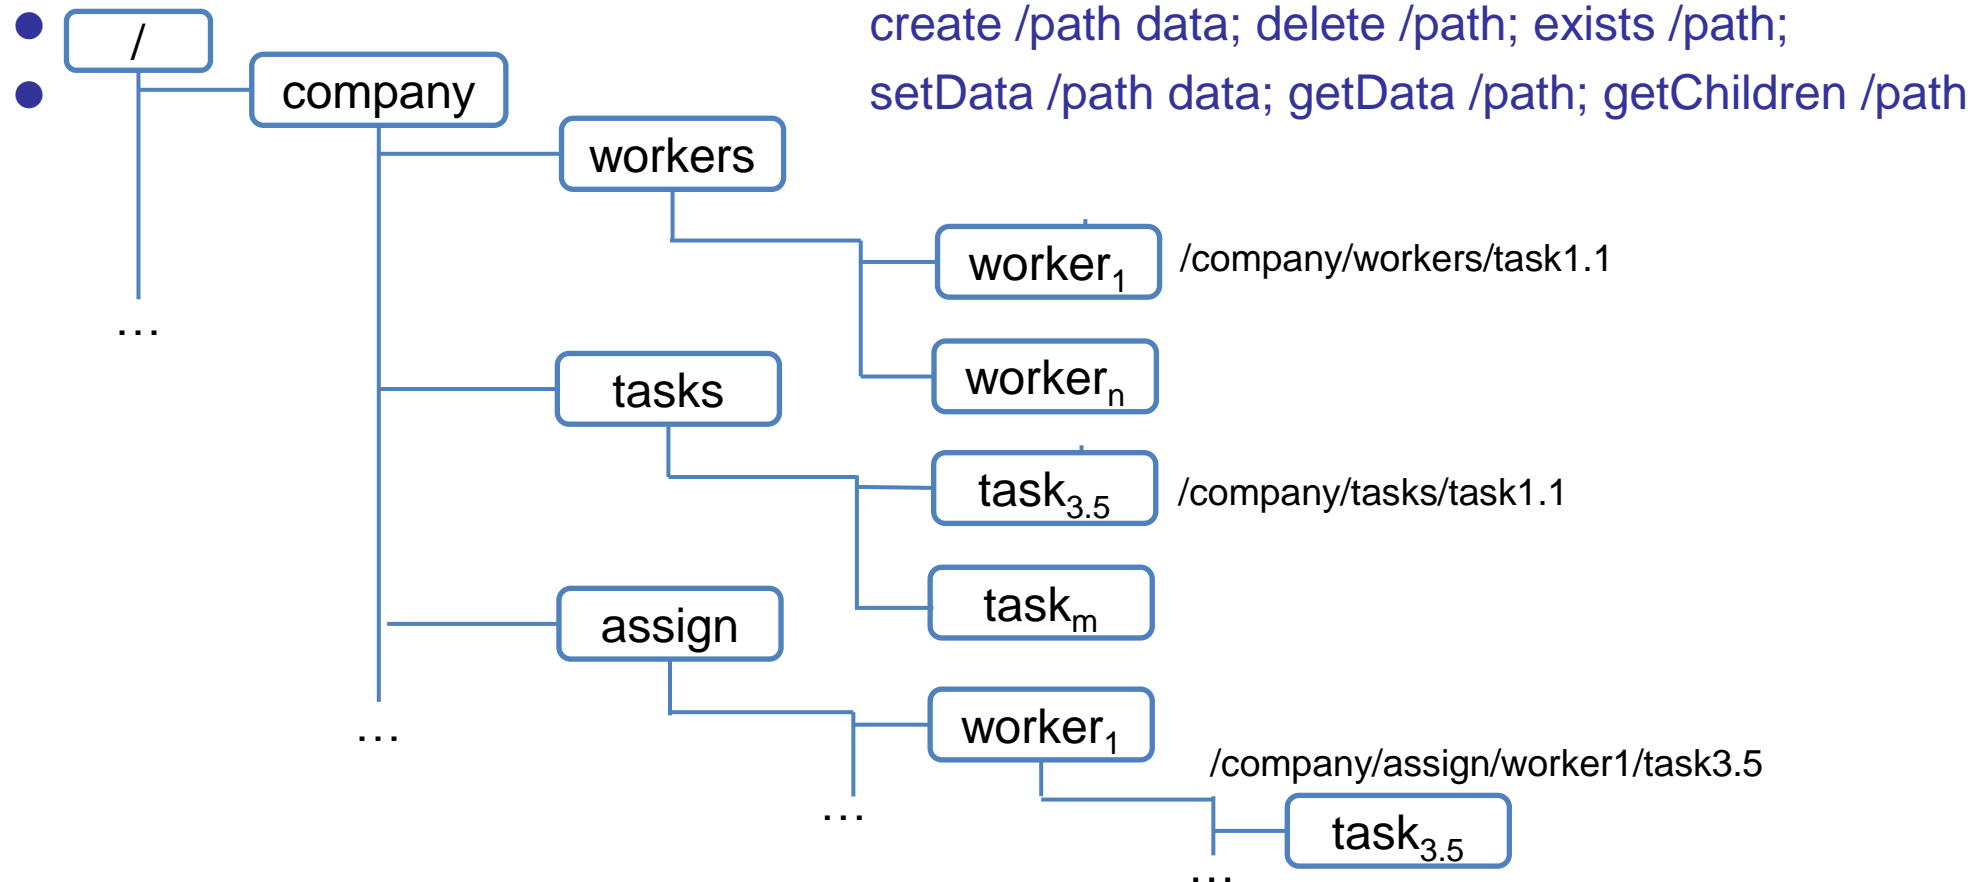
- <https://spring.io/projects/spring-cloud-zookeeper>
 - Service Discovery
 - Client side load-balancer (Ribbon, Spring Cloud Netflix)
 - Ribbon is a Inter Process Communication (remote procedure calls) library with built in software load balancers (<https://github.com/Netflix/ribbon>)
 - Supports Zuul, a gateway service that provides dynamic routing, monitoring, resiliency, security, ... (<https://github.com/Netflix/zuul>)
 - Service Discovery
 - Zookeeper as a data store (e.g., OSGi Remote Services)

■ The Zookeeper system is a reliable node (znode) persistence server.

- A znode is part of a tree structure starting in a root '/', the root znode, and linked to zero or more children (it is a mix of file and a directory)
 - By initialization, Zookeeper starts with a root with one znode, the zookeeper znode that has a child quota



- It can be used to reliably store state through a simple API.



- The zookeeper is an application domain independent system (infrastructure?) responsible for reliably make persistent znodes structured according the needs of the informatics system that uses it as a composite element;
- The clients have the above mentioned six simple operations
 - A znode can't be partially changed (only overridden / substituido)
 - App operations over znodes are idempotent (discussed later)
- A znode can be
 - Persistent (CreateMode.PERSISTENT)
 - Is maintained persistent even when the client that create it disconnects
 - Ephemeral (CreateMode.EPHEMERAL)
 - The znode is discarded when the client that create it disconnects
 - Sequential (Appends a sequence number to a node id)
 - Persistent (CreateMode.PERSISTENT_SEQUENTIAL)
 - Ephemeral CreateMode.EPHEMERAL_SEQUENTIAL)

■ A client accesses zookeeper server based on a Session

- Preserve FIFO order for any sequence of calls (through TCP client configured port)
 - For a same session;
 - if an initial session fails and another is established, the order is not any more guaranteed
- A session is created through a ZooKeeper object
 - new ZooKeeper (ZOOKEEPER_SERVER, TimeOut [, Watcher]); where ZOOKEEPER_SERVER is the hostname or IpAddress of the host where Zookeeper server is running (or the Ensemble list of servers; a list host:clientPort)
- A client session can also be created by starting the zkCli command
 - The prompt is [zk: siserver0.local(CONNECTED) 0] help
 - Through the help command is possible to access available commands.
- ACL (Access Control List) Permissions
 - ZooKeeper supports the following permissions:
 - CREATE: you can create a child node
 - READ: you can get data from a node and list its children.
 - WRITE: you can set data for a node
 - DELETE: you can delete a child node
 - ADMIN: you can set permissions

■ According the Consistency, Availability, and Partition-tolerance (CAP) theorem

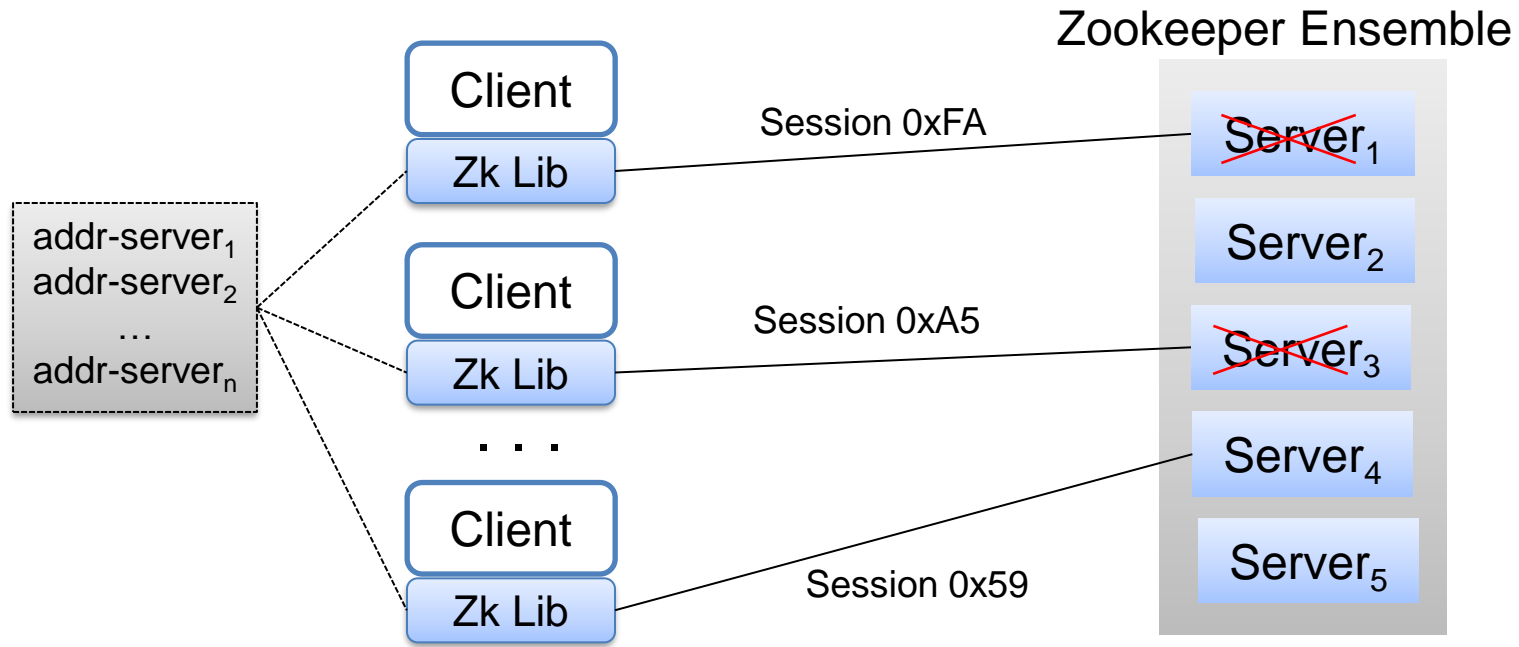
- No one distributed system can handle the three dimensions simultaneously; by Eric Brewer, “*Towards robust distributed systems*”, 2000 ([link](#))
- Later revisited “CAP twelve years later- Discusses further the theorem”, 2012 ([link](#))

■ The Zookeeper system challenges such a possibility through a server, the zookeeper server, that can be configured considering the following modes:

- Standalone (not reliable) (used in OSGi remote services)
- Quorum, based on an **Ensemble** of servers
 - That, depending on the number of servers, failure of some is made transparent to client systems
 - The minimal number of ensemble servers is established by the minimum quorum that requires a majority of servers, like in real life general assemblies

■ The reliability is related to its implementation strategy

- Reliability, or at least an higher reliability level, requires the activation of the quorum mode
- The quorum mode requires a group of servers to guarantee some level of replication of zookeep server state
 - Such a set of servers is named **Ensemble**



■ The Quorum concept

- The quorum is in fact the minimal number of servers necessary for Zookeeper to work (to guarantee znode state replication)
 - Similar to the minimal number of presences in a general assembly to be able to take decisions, as mentioned already
- The question now is, what should be the minimal number of servers to make a quorum?
 - Suppose in an Ensemble of 5 we assume a quorum of 2, is it enough?
 - Suppose that servers 1 and 3 acknowledge the creation of the node /osgi, but before the change is propagated to the other servers (2, 4, and 5) the servers 1 and 3 split the partition (known as **Split-brain** situation), by some failure.
 - Based on the established minimal quorum of 2, client continues convinced that /osgi was persisted. However as the partition happened before the replication process was complete, in fact, servers 2, 4, and 5 do not have the created znode
- Important to consider
 - A majority of nodes is necessary to guarantee consistency of the replicated state (FSM, our vector or other critical data)
 - How can we scale a stateful service element?
 - How to deal with scale?

■ The Quorum concept

- A quorum requires a majority of servers of an ensemble
 - This means that for an ensemble of 5 the quorum is at least 3
 - In this case the tolerance to failure is of two servers
- Formally if E represents an ensemble with $|E|$ servers and N_F the number of faulty servers,
 - From the assumption of Lamport, $N > 3N_F$ consensus can be obtained, and
 - Therefore the minimal quorum for one faulty server of an Ensemble, is
 - $N = 3 \cdot N_F + 1$; for the minimal quorum of 3, $|E| = 3 \cdot 1 + 1 = 4$, the required number of nodes (independent servers)

■ Noteworthy

- The Service elements instances shall be running on failure-independent execution environments
 - E.g., each node running on different racks with independent power supplying
- If disaster fault tolerance is to consider
 - Nodes need to be in different geographical locations, and communication infrastructure needs to show redundancy

■ Coordination Strategy

- Based on the Zookeeper Atomic Broadcast (Zab) protocol
 - It was proposed by Flavio P. Junqueira, Benjamin C. Reed, and Marco Serafini, from Yahoo! Research, in the paper “Zab: High-performance broadcast for primary-backup systems”, 2011 ([link](#))
 - It is an evolution/simplification from the Paxos consensus algorithm initially proposed by Leslie Lamport in the paper “The Part-Time Parliament”, 2000 ([link](#))
 - “Recent archaeological discoveries on the island of **Paxos** reveal that the **parliament functioned** despite the peripatetic propensity of its **part-time legislators**. The legislators maintained **consistent copies of the parliamentary record**, despite their frequent forays from the chamber and the forgetfulness of their messengers. The Paxos parliament’s protocol provides a new way of implementing the **state-machine approach to the design of distributed systems**.”
- It is based also on the Byzantine Generals problem
 - How to guarantee that only **loyal generals** will **reach agreement** on a **military strategy**?

■ Zookeeper

- Considers three classes of servers:
 - Leader, Followers, and Observers

■ The main role of each class of servers

- Leader
 - The one that **coordinates commits** (of zxid transactions) of the server's state
 - The **state change** is processes by the **leader**
- Followers
 - **Forward state changes** to the Leader
 - Process **reads**
 - Initiate **leader election** if they can't contact the leader
- Observers
 - Process reads (scalability reasons)
 - As they don't count for quorum, their number can vary according the need to solve bottlenecks for reads making more Observers to help and free Followers for updates
 - Do not participate in the election process

■ Operations over znodes

- Read
 - exists /path; getData /path; getChildren /path
- Update
 - create /path data; delete /path; setData /path data

■ Transactions

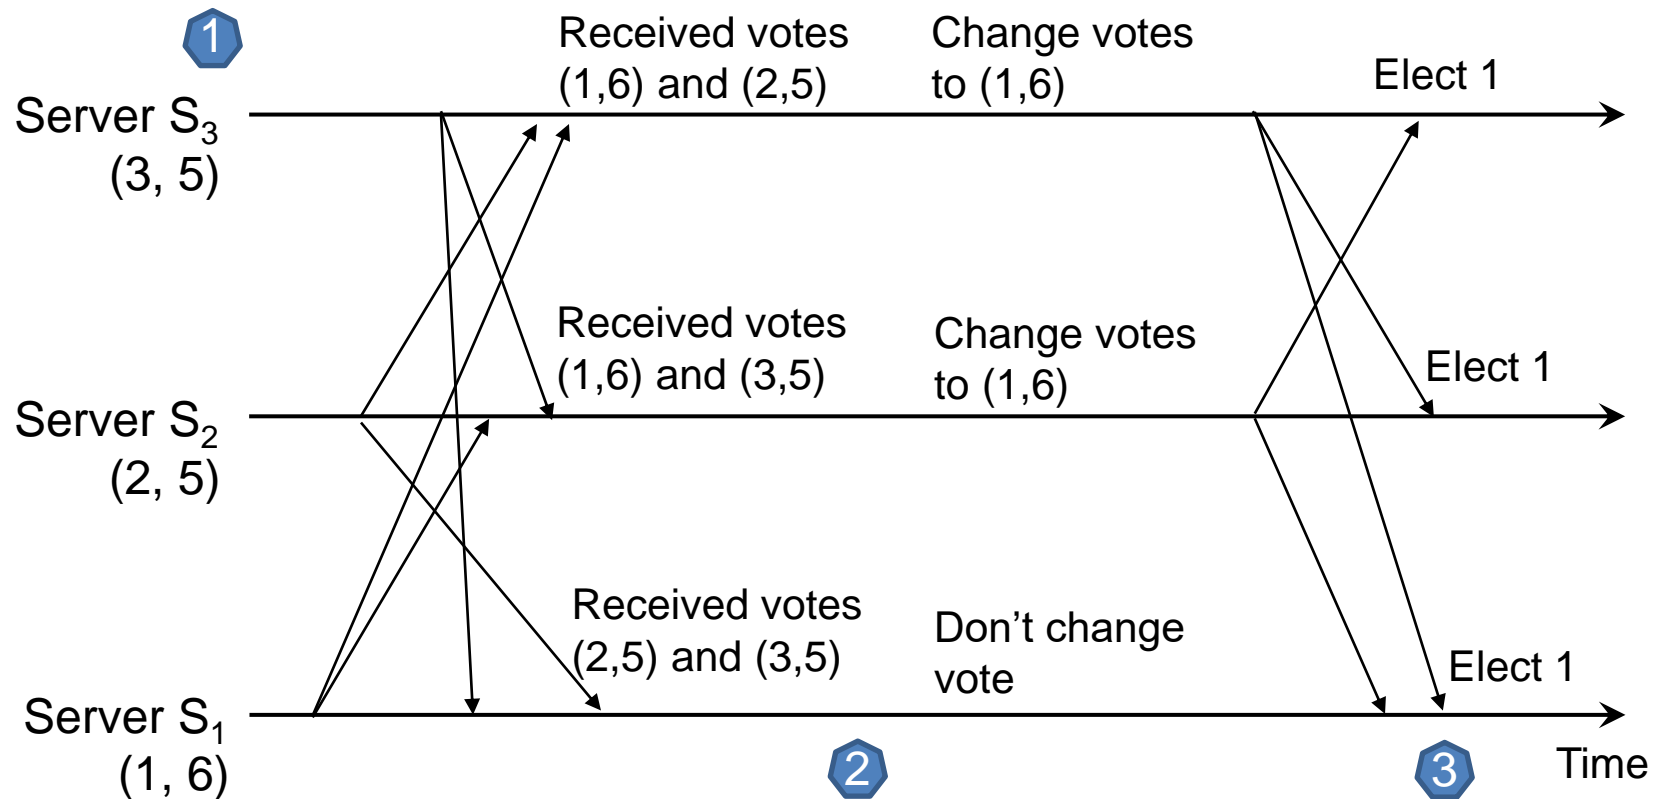
- The transactions are **idempotent** operations, i.e., operations can be replayed without changing the expected result
- A zookeeper transaction, named zxid is a 64 bit long with two parts
 - Epoch, 32 bits
 - Counter, 32 bits

■ Leader Election

- Is the one chosen among ensemble members to take the Leader role
- A main purpose is to order client requests to change the zookeeper state (create /path data; delete /path; setData /path data)
- Each server starts in the LOOKING state
 - Elect a new server, or
 - Find the existing one
- If election is needed, depending on the result
 - One server enters the LEADING state, and
 - The others enter the FOLLOWING state

■ Leader Election Algorithm

- Leader election messages
 - ***leader election notifications***, or ***notifications***
- When a server enters the LOOKING state it sends a batch of notification messages
 - The message contains its current vote, the server id (sid) and the most recent executed transaction (zxid)
- Upon receiving a vote, a server changes its vote according the following rules:
 - Let $\text{vote}(\text{Id}, \text{Zxid})$ be the identifier and the zxid in the current received vote, whereas myZxid and mySid are the values of the receiver itself
 - If $(\text{voteZxid} < \text{myZxid})$ or $(\text{voteZxid} = \text{myZxid} \text{ and } \text{voteId} > \text{mySid})$, keep the current vote
 - Otherwise, change my vote by assigning myZxid to voteZxid and mySid to voteZxid
- The server that is most up to date is the winner
 - I.e., has an higher value for the zxid
 - For a same zxid, prevails the higher Sid server

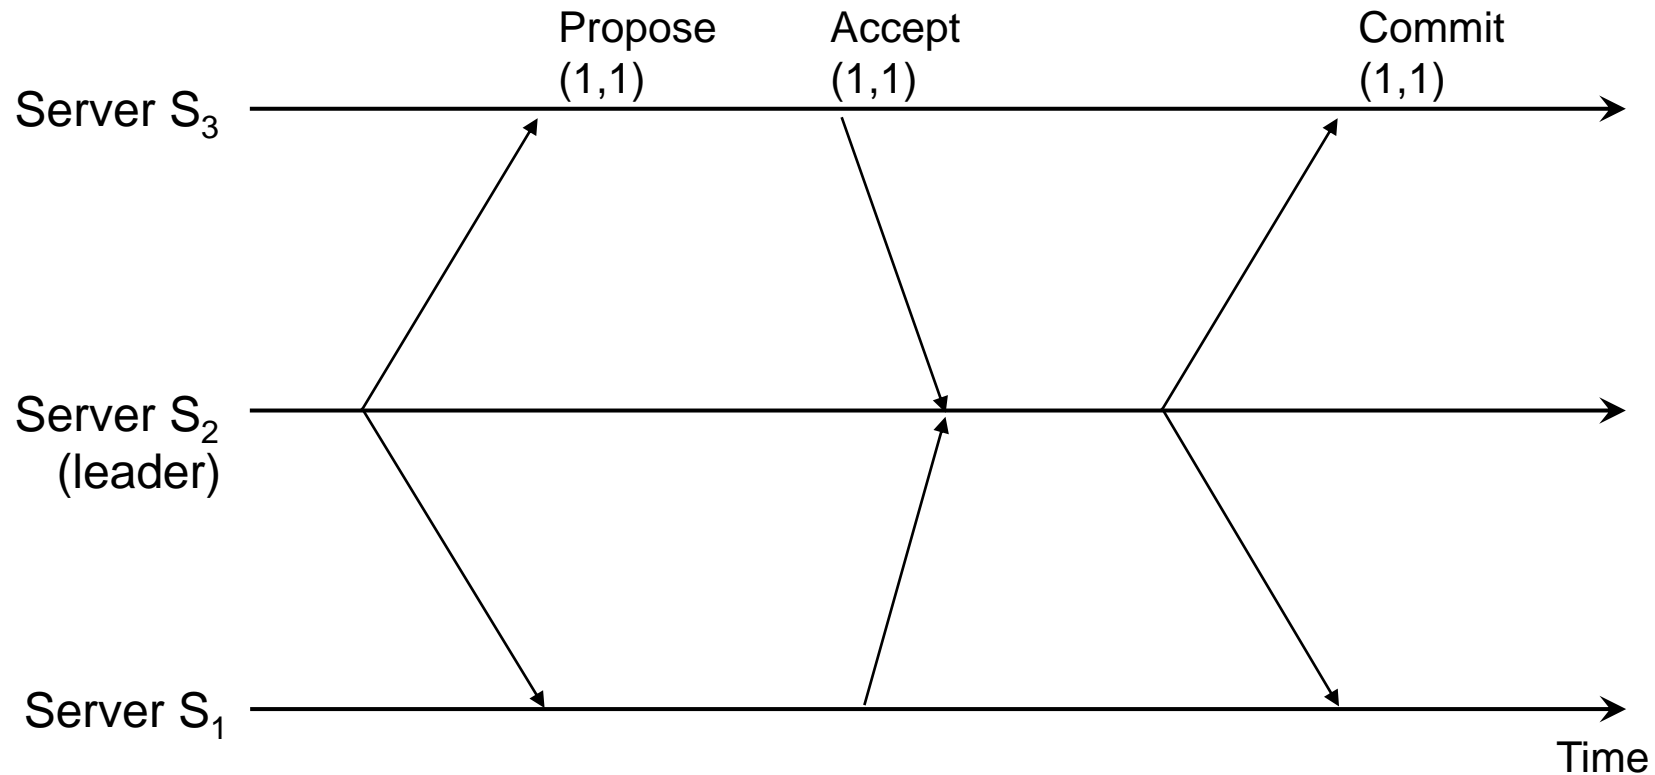


- 1 Server S_1 starts with vote (1, 6), server S_2 with (2, 5) and S_3 with (3, 5)
- 2 Server S_1 and S_3 change their vote to (1, 6) and send a new batch of notifications
- 3 All the three servers receive the same vote from a quorum and elect server S_1

■ The Zab (Zookeeper Atomic Broadcast) protocol

- The question now is how to determine the commit of a transaction
- It refers any change to the Zookeeper DataTree, the tree of znodes
 - Any change, corresponds to a DataTree state change that needs to be maintained consistent in the quorum of servers
- The protocol to commit a transaction resembles the two-phase commit algorithms as established by the X/Open standard
 - 1. The Leader sends a PROPOSAL message p to all the followers
 - 2. Upon receiving p, followers respond to the leader with an ACK, informing the leader that it has accepted the proposal
 - 3. Upon receiving acknowledgments from a quorum (including the leader), the leader sends a message informing the followers to COMMIT the transaction
- Zab guarantees that:
 - If the leader broadcast T and T' in that order, each server must commit T before T'
 - If any server commits transactions T and T' in that order, all servers must commit T before T'

Commit Transactions



- 1 Leader sends propose messages
- 2 Followers ack the proposal
- 3 Leader commit the proposal

From: Zookeeper, Distributed Process Coordination;
ISBN: 9781449361303, O'Reilly, 2013

Initialize zookeeper server in standalone mode

■ Install from Apache zookeeper project ([stable version](#))

- Installation: C:\Java\apache-zookeeper-3.9.2-bin\
 - On Linux (Ubuntu) /opt/zookeeper
- C:\Java\apache-zookeeper-3.9.2-bin\conf\zoo.cfg
 - tickTime=2000
 - dataDir=\tmp\zookeeper
 - On Linux: /tmp/zookeeper
 - clientPort=2181
 - **tickTime** - time unit in milliseconds (heartbeats, minimum session timeout $2 * tickTime$)
 - **dataDir** - in-memory database and transaction log
 - **clientPort** - the port to listen for client connections
- Start the server
 - C:\...>Java\apache-zookeeper-3.9.2-bin\bin\zkServer.cmd start
- Start a client
 - C:\...>\ Java\apache-zookeeper-3.9.2-bin \bin\zkCli.cmd –server siserver0.local:2181

■ Install Zookeeper in all the participating servers

- C:\Java\apache-zookeeper-3.9.2-bin\conf\zoo.cfg
 - tickTime=2000
 - dataDir=/tmp/zookeeper
 - clientPort=2181
 - server.1=siserver0.local:2888:3888
 - server.2=siserver1.local:2888:3888
 - server.3=siserver2.local:2888:3888
- *Replicate the configuration file to all the servers*
- Start the servers
 - si@siserver0:~\$ sudo /usr/local/zookeeper/bin/zkServer.sh start
 - si@siserver1:~\$ sudo /usr/local/zookeeper/bin/zkServer.sh start
 - si@siserver2:~\$ sudo /usr/local/zookeeper/bin/zkServer.sh start
- Start a client
 - C:\...>\Java\apache-zookeeper-3.9.2-bin\bin\zkCli.cmd –server siserver0.local:2181
 - C:\...>\ Java\apache-zookeeper-3.9.2-bin\bin\zkCli.cmd –server "siserver0.local:2181,siserver1.local:2181,siserver2.local:2181"

■ The technology Landscape of an Organization (IT)

- Is composed of (ISoS objects)
 - One ISystem_0 ,
 - Zero or more ISystem
- An ISystem (Informatics System), is composed of
 - One or more Cooperation Enable Service (CES)
- A CES is composed of
 - One or more Service, as independent (autonomous) computational entity
 - The Service entity is in fact the artifact that embeds computational logic and resources necessary to operationalize a computational responsibility.

■ The ISystem_0 serves as $\text{ISystem}/\text{CES}/\text{Service}$ Registry

- A REST interface (I_0) accessible at
 - `isos.<organization_domain>:2058`
- Has an Administration user interface (Admin)
 - Navigation through the ISoS tree ($\text{ISystem}/\text{CES}/\text{Service}$)

Reliable ISoS Reference Implementation

