
Transactions – *locks* the **ACID** properties

- Atomicity
- Consistency
- Isolation
- Durability

■ Exclusive access to data elements (critical resource) - lock

- Each transaction promotes exclusive access to data elements

■ Depending on the access mode, the locks can be:

- Read locks (also known as Shared locks)
- Write locks (also known as Exclusive locks)

■ Access to a data element

- Before a read, the transaction sets a ReadLock
- Before a write, the transaction sets a WriteLock

■ Conflicts

- A ReadLocks establishes a conflict with WriteLocks
- WriteLocks establishes conflict with ReadLocks and WriteLocks

■ Obtaining the locks to be associated to a transaction

- Only if no conflict locks in relation to a same data element;
 - Can obtain a ReadLock to access X (data element) if no other transaction has a WriteLock on X (data element)
 - Can obtain a WriteLock on X (data element) if no other transaction has a ReadLock or a WriteLock on X

■ Impacts on performance (parallelism)

■ A lock set by a transaction

- Delays other transactions that establishes conflict locks

■ Deadlocks

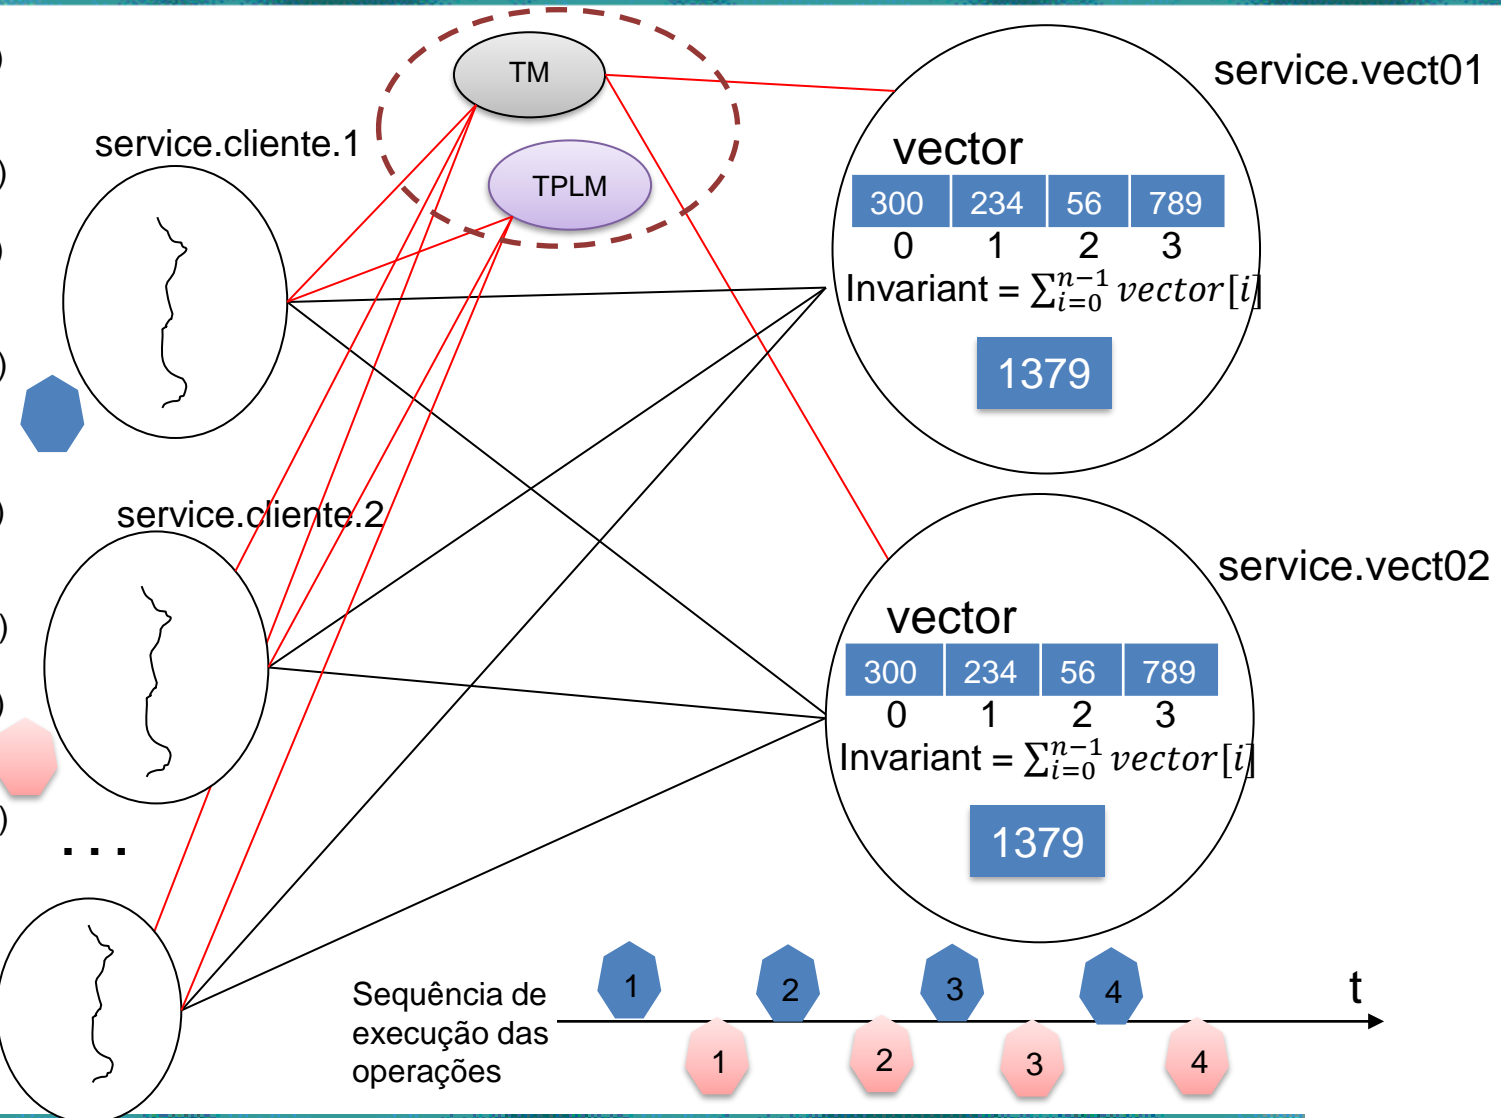
- Requires strategies to define data model structure, and coordination policies, being fundamental aspects
 - Locking “granularity”

Serviços concorrentes no acesso a múltiplos ser.vectxx

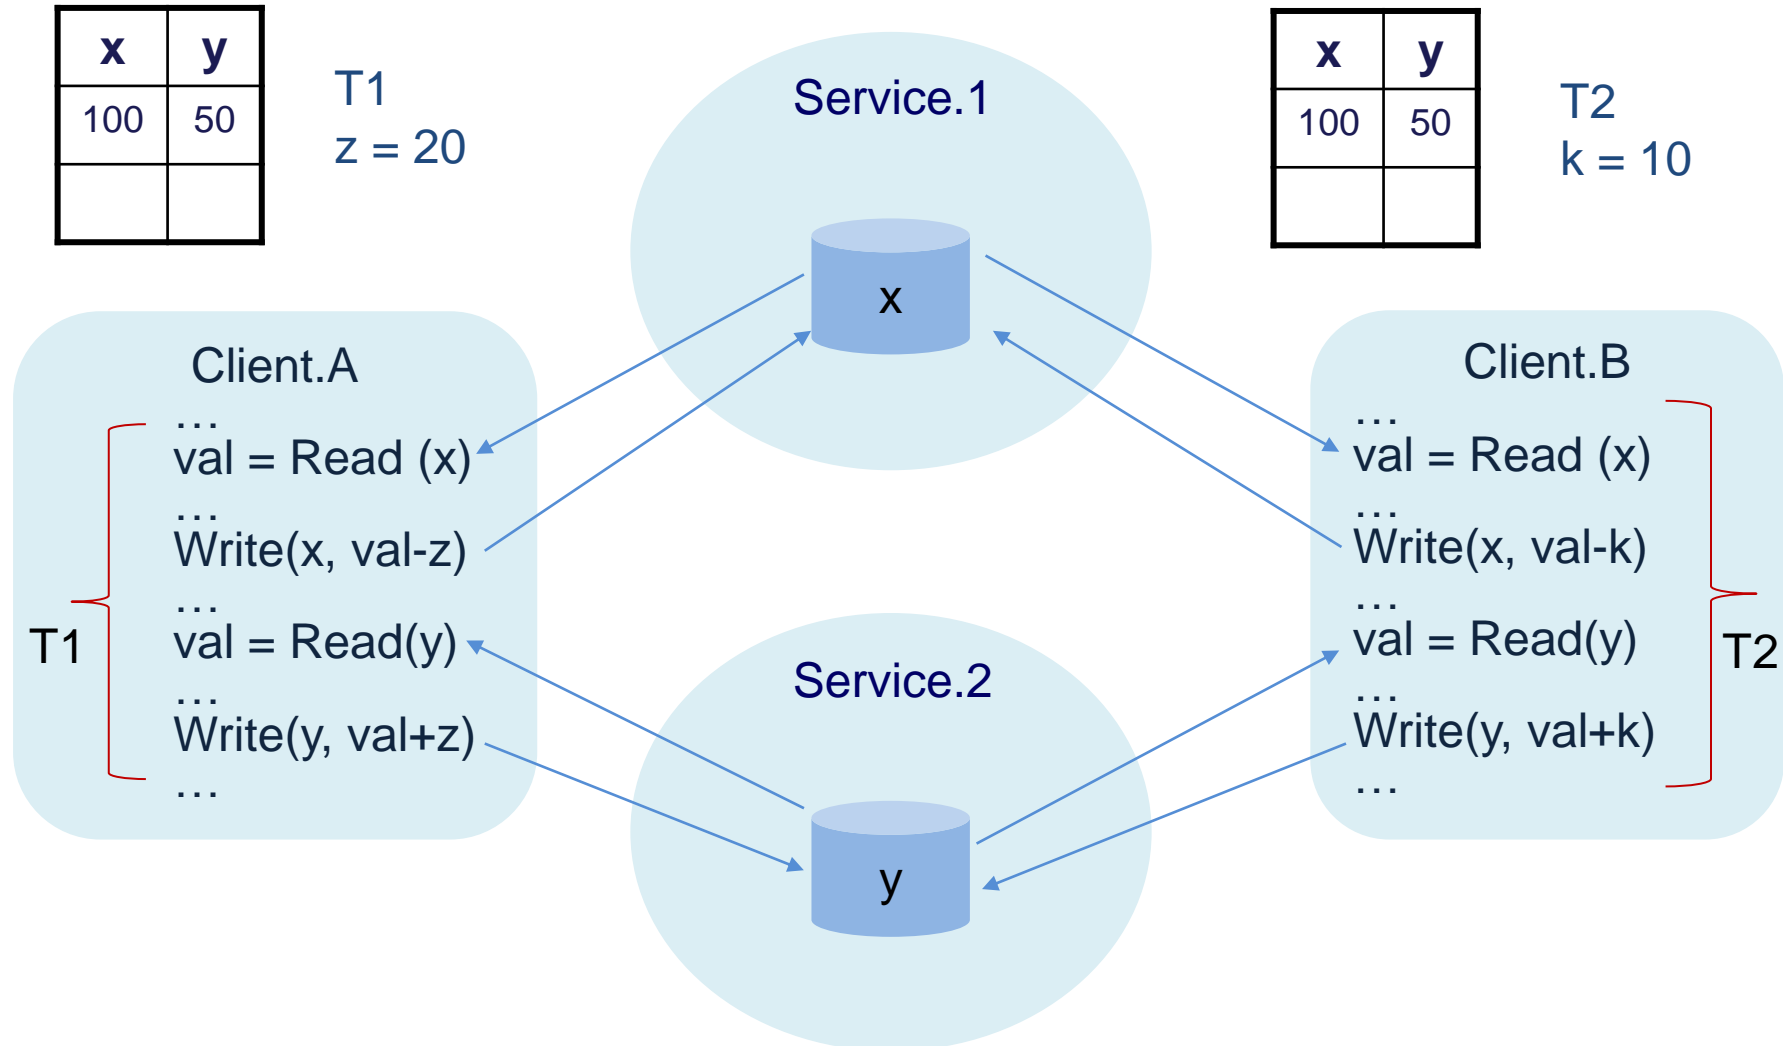
1. $v = \text{read}(\text{vect01.vect}[0])$
 $\text{res} = v - x // x = 100$
Espera(10 segundos)
 2. $\text{write}(\text{vect01.vec}[0], \text{res})$
Espera(10 segundos)
 3. $v = \text{read}(\text{vect02.vect}[2])$
 $\text{res} = v + x$
Espera(10 segundos)
 4. $\text{write}(\text{vect02.vec}[2], \text{res})$

1. $v = \text{read}(\text{vect01.vect}[0])$
 $\text{res} = v - x // x = 100$
Espera(10 segundos)
 2. $\text{write}(\text{vect01.vec}[0], \text{res})$
Espera(10 segundos)
 3. $v = \text{read}(\text{vect02.vect}[2])$
 $\text{res} = v + x$
Espera(10 segundos)
 4. $\text{write}(\text{vect02.vec}[2], \text{res})$

service.cliente.n



Scenario with two data elements, x and y



Concurrent execution of T1 and T2 on x and y elements

(Concurrent execution of T_1 and T_2) E =

$rl_1[x] \ r_1[x] \ ru_1[x]$
T₁ reads x ...

$rl_2[y] \ r_2[y] \ ru_2[y] \ rl_2[x]$
 $r_2[x] \ ru_2[x] \ wl_2[y]$
 $w_2[y+k] \ wu_2[y] \ wl_2[x]$
 $w_2[x-k] \ wu_2[x]$
T₂ reads and writes y and x

$wl_1[y] \ r_1[y] \ w_1[y+z] \ wu_1[y]$
 $wl_1[x] \ w_1[x-z] \ wu_1[x]$
...T₁ reads y and writes y and x

Serialized execution?

Demonstrating the violation of the Serializable rule

Removing the locks and unlock from the execution of E we obtain E'

$E' = r_1[x] \parallel r_2[y] \parallel r_2[x] \parallel w_2[x-k] \parallel w_2[y+k] \parallel r_1[y] \parallel w_1[y+z] \parallel w_1[x-z]$

Serialized execution?

x	y
100	50

Invariant = $x+y$ (150);

T1 $z = 20$

T2 $k = 10$

x	y
100	50

x	y
100	50

$T_1 T_2 = r_1[x] \parallel r_1[y] \parallel w_1[y+z] \parallel w_1[x-z] \parallel r_2[y] \parallel r_2[x] \parallel w_2[x-k] \parallel w_2[y+k]$

The two possible serialized executions:

$T_2 T_1 = r_2[y] \parallel r_2[x] \parallel w_2[x-k] \parallel w_2[y+k] \parallel r_1[x] \parallel w_1[x-z] \parallel r_1[y] \parallel w_1[y+z]$

- The Transaction T_1 unlocks x before obtaining a write lock on y
- T_1 opens an opportunity window for other concurrent transaction to meanwhile access the data element

■ A transaction shall obtain all the locks before release them

- Founded on a theorem demonstrated by Ullman (1982)
 - Principles of Transaction Processing, by Philip A. Bernstein - Chap. 6

■ In the previous example:

- T1 shouldn't release the lock on x so early (and it should be a write lock)
- As there is a pending update (the K value) $wu_1[x]$ obtained at the beginning should be released only at the end

■ Important architectural decisions

- The granularity of the data elements
- The need for a LockManager to guarantee the consistency during transactions execution (concurrency coordinator)

■ Operations (methods)

- getLocks(transactionID, lockElements)
- releaseLocks(transactionID, lockElements)
- unlock(transactionID) // release all locks associated to the transaction

■ LockElements[]

- LockElement
 - DataItem, LockMode
 - lockMode = read | write

■ Questions

- How to implement it on the Jini framework?

■ Evaluation scenario

- Transferences between accounts from different bank (services), based on low level read and write operations.

Alternative Concurrency Coordination Strategies

■ Optimistic concurrency control (OCC)

- *“Optimistic concurrency control (OCC) ..., operates on the principle that multiple transactions can frequently complete without interfering with each other.” [\[ref\]](#).*

■ Multiversion concurrency control (MVCC)

- *“MVCC (Multi-Version Concurrency Control) is a way to improve database concurrency by storing multiple row versions as they have been at different points in time.” [\[ref\]](#)*

Long lived transactions (LLTs) - SAGA

- Molina defined in 1987 [\[ref\]](#) the concept of Long-lived transactions (LLTs) and suggested the name of **Saga** when
 - *“A LLT is a saga if it can be written as a sequence of transactions that can be interleaved with other transactions T ”*
- The concept was meanwhile adopted in the context of microservices [\[ref\]](#) considering the case of long transactions, e.g., when the idea of Roll-back is not applicable because a failure requires the execution of a process to return to the prior state.
 - An example of organizing a trip where reservation services, e.g., flight, hotel, transfer, require hours or days to execute, and if necessary to cancel, there is commonly a reimbursement payment associated that needs to be processed.
- The term **Long Running Transactions** is also applied to web services [\[ref\]](#), that this paper associate to Saga concept.
 - *“With sagas, many small atomic transactions are wrapped by a larger longer running transaction. Each small atomic transaction is paired with a **compensation** handler that is capable of reversing the activity done in the atomic transaction.”*

